

HTTP/1.1 vs HTTP/2

a performance analysis

Ragnar Lönn, Load Impact

ragnar@loadimpact.com

@ragnarlonn

Load...Impact?

- **LOAD TESTING SAAS**
- **LAUNCHED 2009**
- **150,000+ ACCOUNTS**
- **1.5M+ LOAD TESTS**
- **APPLICATION-POSITIVE**




Part 1: The protocol

http:




What is the problem?



Source: Mike Belshe, Google


Load time and delay




Source: Mike Belshe, Google

Some history


1996

- **Web pages were tiny**
 - **Internet bandwidth was tinier**
 - **HTTP/1.0 spec released (RFC1945)**
 - **Work started on HTTP/1.1**
- 
- A decorative graphic in the bottom right corner consisting of several overlapping circles in shades of light green and grey, creating a stylized, organic shape.


Page load circa 1996



Page composition



Page load today



Some more history

1997

- HTTP/1.1 spec released (RFC2068)
- A lot of stuff already implemented...
- Some performance improvements

Pipelining


Connection
re-use

Improved
caching


Range
requests




Connection re-use



Moar connections!



Pipelining - if it worked



Performance optimizations

a.k.a. “ugly hacks”

- **Spriting**
- Inlining
- Concatenating
- Sharding



Spriting

Performance optimizations

a.k.a. “ugly hacks”


- Spriting
- Inlining
- Concatenating
- Sharding

```
GET host1.domain.com/resourceA  
GET host2.domain.com/resourceB  
GET host3.domain.com/resourceC  
...
```


Sharding

Ugh.

Yes. So what if there was a protocol that:

- **...was less sensitive to network delay**
 - **...fixed pipelining and the head of line blocking**
 - **...performed well regardless of the number of TCP connections used**
 - **...used the same addressing/URL schemes etc for compatibility with HTTP 1.1**
- 

HTTP/2 multiplexing



Multiplexing



Multiplexing



HTTP/2 at a glance

- HTTP/2 is binary and (mostly) encrypted
- “Upgrade:” header or TLS ALPN negotiation
- Multiplexing is done with “streams”
 - Streams can be prioritized, re-prioritized and cancelled at any time
 - Streams can have dependencies
 - Streams have individual flow control
- Headers are compressed
- Servers can push content to clients



HTTP/2 penetration

- >70% of installed browsers support HTTP/2



HTTP/2 server side

facebook

Google twitter



Microsoft
IIS



CADDY



APACHE
HTTP SERVER

traffic:server™

LITESPEED

NGINX

HTTP/2 penetration

- **24% of Firefox traffic is HTTP/2**
- **18% of Alexa top 500 support HTTP/2**



Part 2: The experiment

Objective:

Try to get a sense of real-world performance impact of going from HTTP/1.1 to HTTP/2



Approach

- Choose a well-known site
- Download all resources used by main page
- Host everything locally, in controlled environment
- Measure load times while simulating different network characteristics



Site: www.amazon.com



- ~230-240 resources to get to onload()
- ...but “only” ~10 javascript files
- ~10 unique source hosts
- ~7 Mbyte data

The setup



Chrome



Linux 3.19.0-25 (Ubuntu)
2 CPU cores, 2G RAM


eth0:1
...
...
eth0:n

- Nginx 1.9.5/http2
- Shimmer Cat 0.1
- Linux Netem

VMware Fusion

MacOS X 10.10.3 (4 CPU cores, 16G RAM)

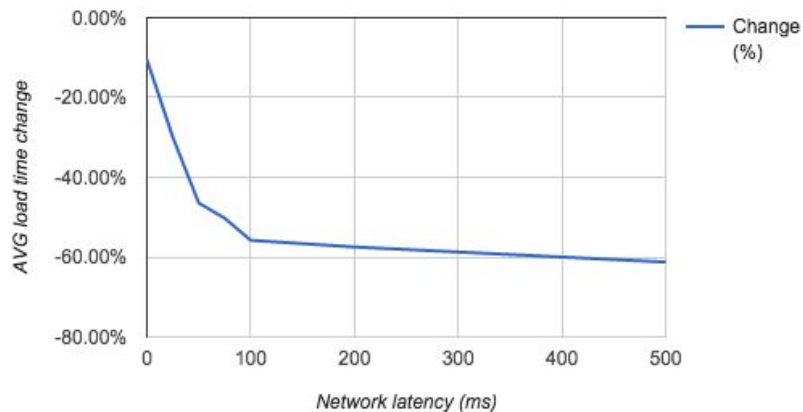
Test results



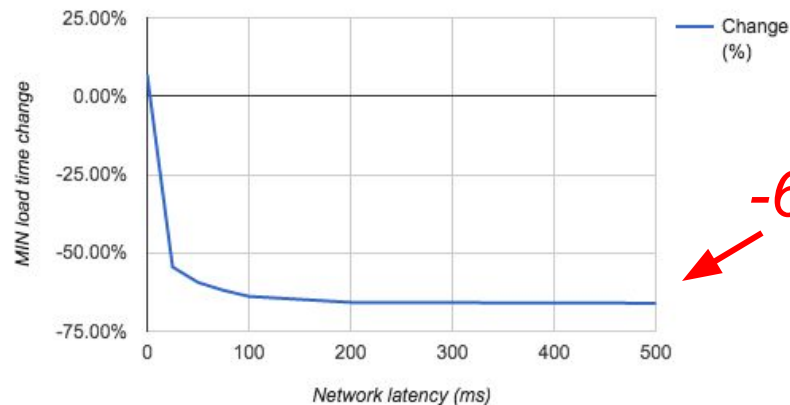
Test results

50-70% reduced load time

HTTP/1.1 - HTTP/2 AVG load time change



HTTP/1.1 - HTTP/2 MIN load time change



But remember:


- HTTP/2 implementations still in their infancy
- Sites are optimized for HTTP/1.1
- Our lab setup is not 100% realistic

So who knows... But still, hey - 60%!



http2.loadimpact.com

Coming soon



Code @ <https://github.com/loadimpact>