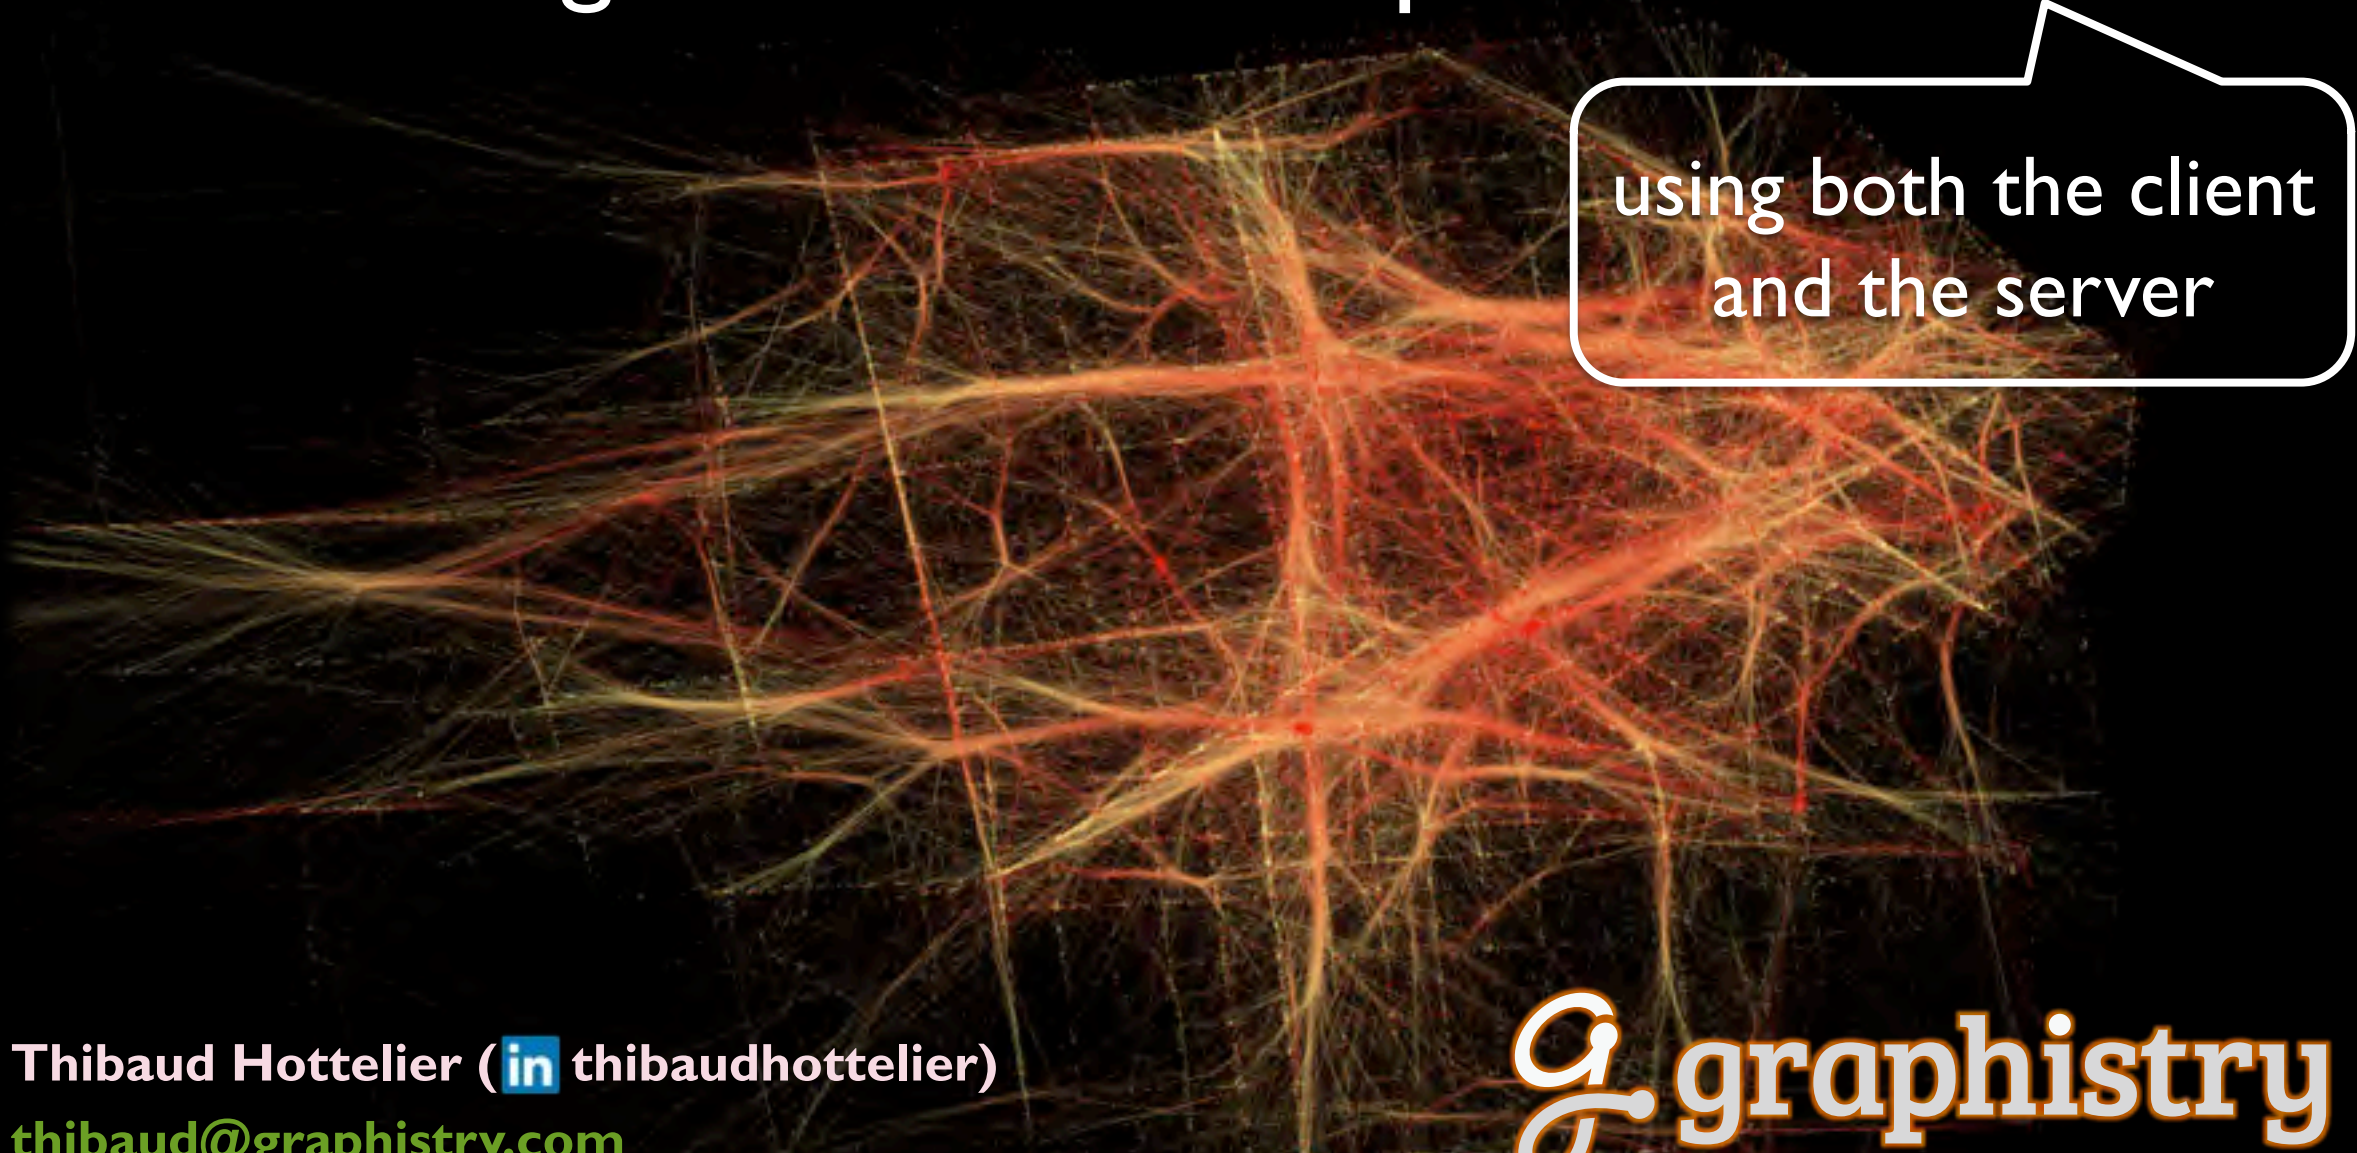


Visualizing Millions of Datapoints with GPUs



using both the client
and the server

Thibaud Hottelier ([in](#) thibaudhottelier)
thibaud@graphistry.com



GRID + SHIELD

 **NVIDIA.**
GEFORCE NOW
THE NEW WAY TO GAME.

\$7.99/MONTH
FIRST 3 MONTHS FREE!*

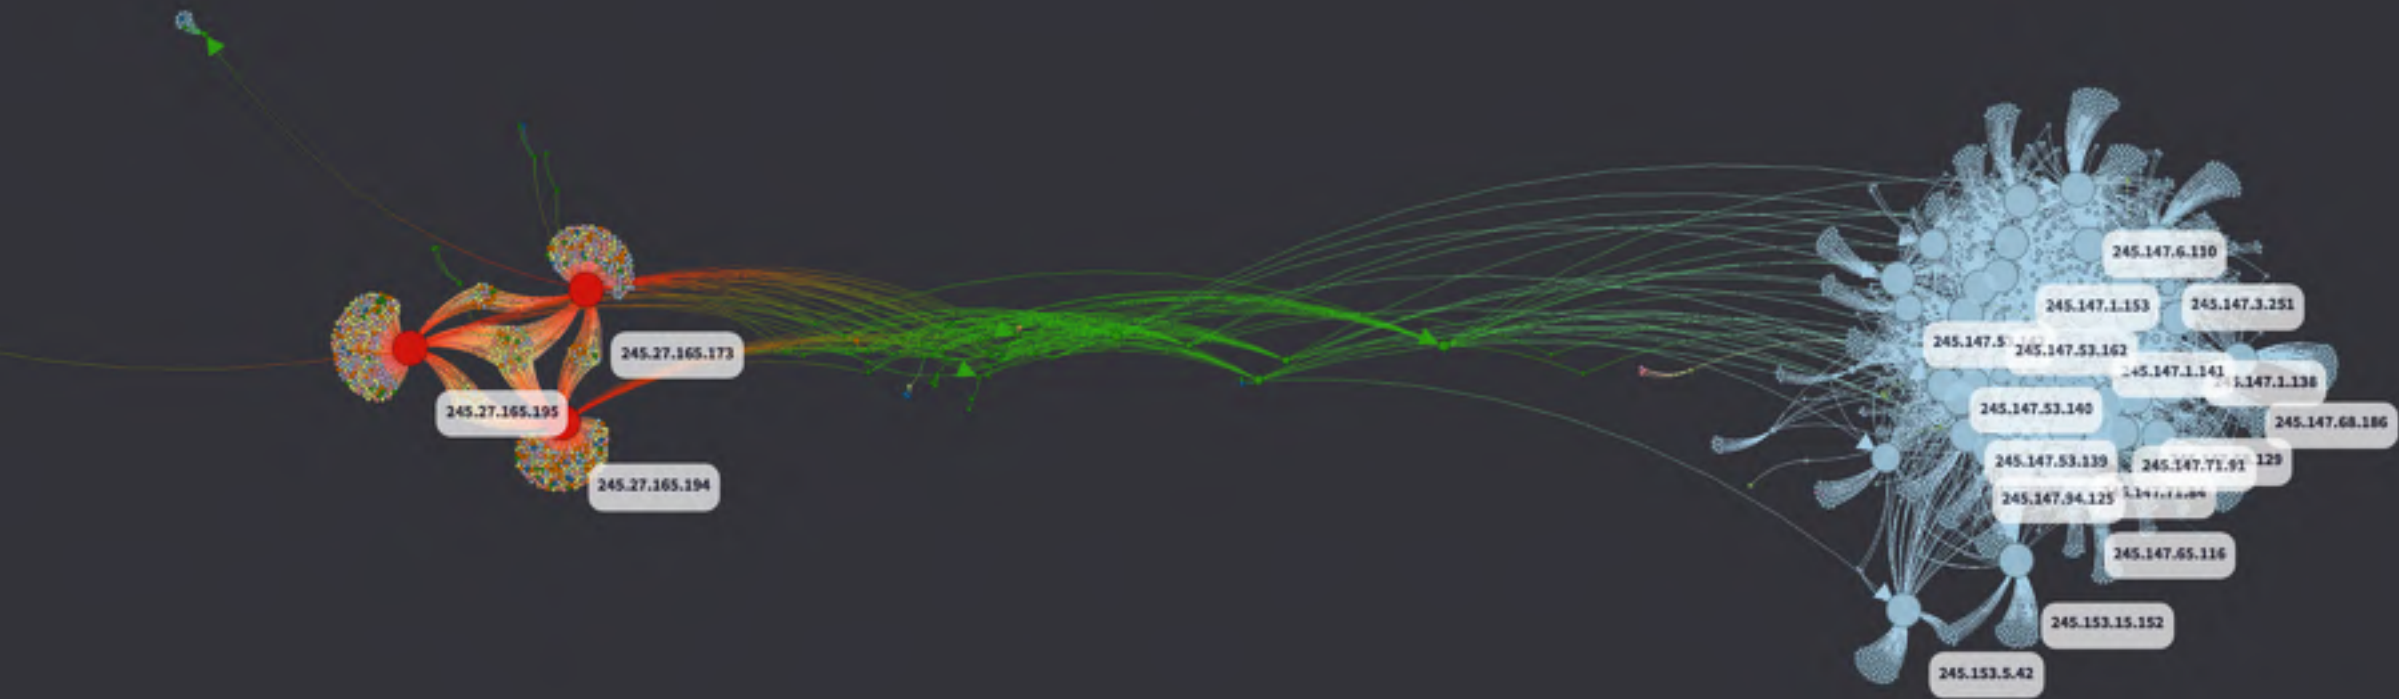
[LEARN MORE](#)

GPUs for Live VR Streaming



GPUs for Machine Learning

GPUs to Map Cyber Attacks



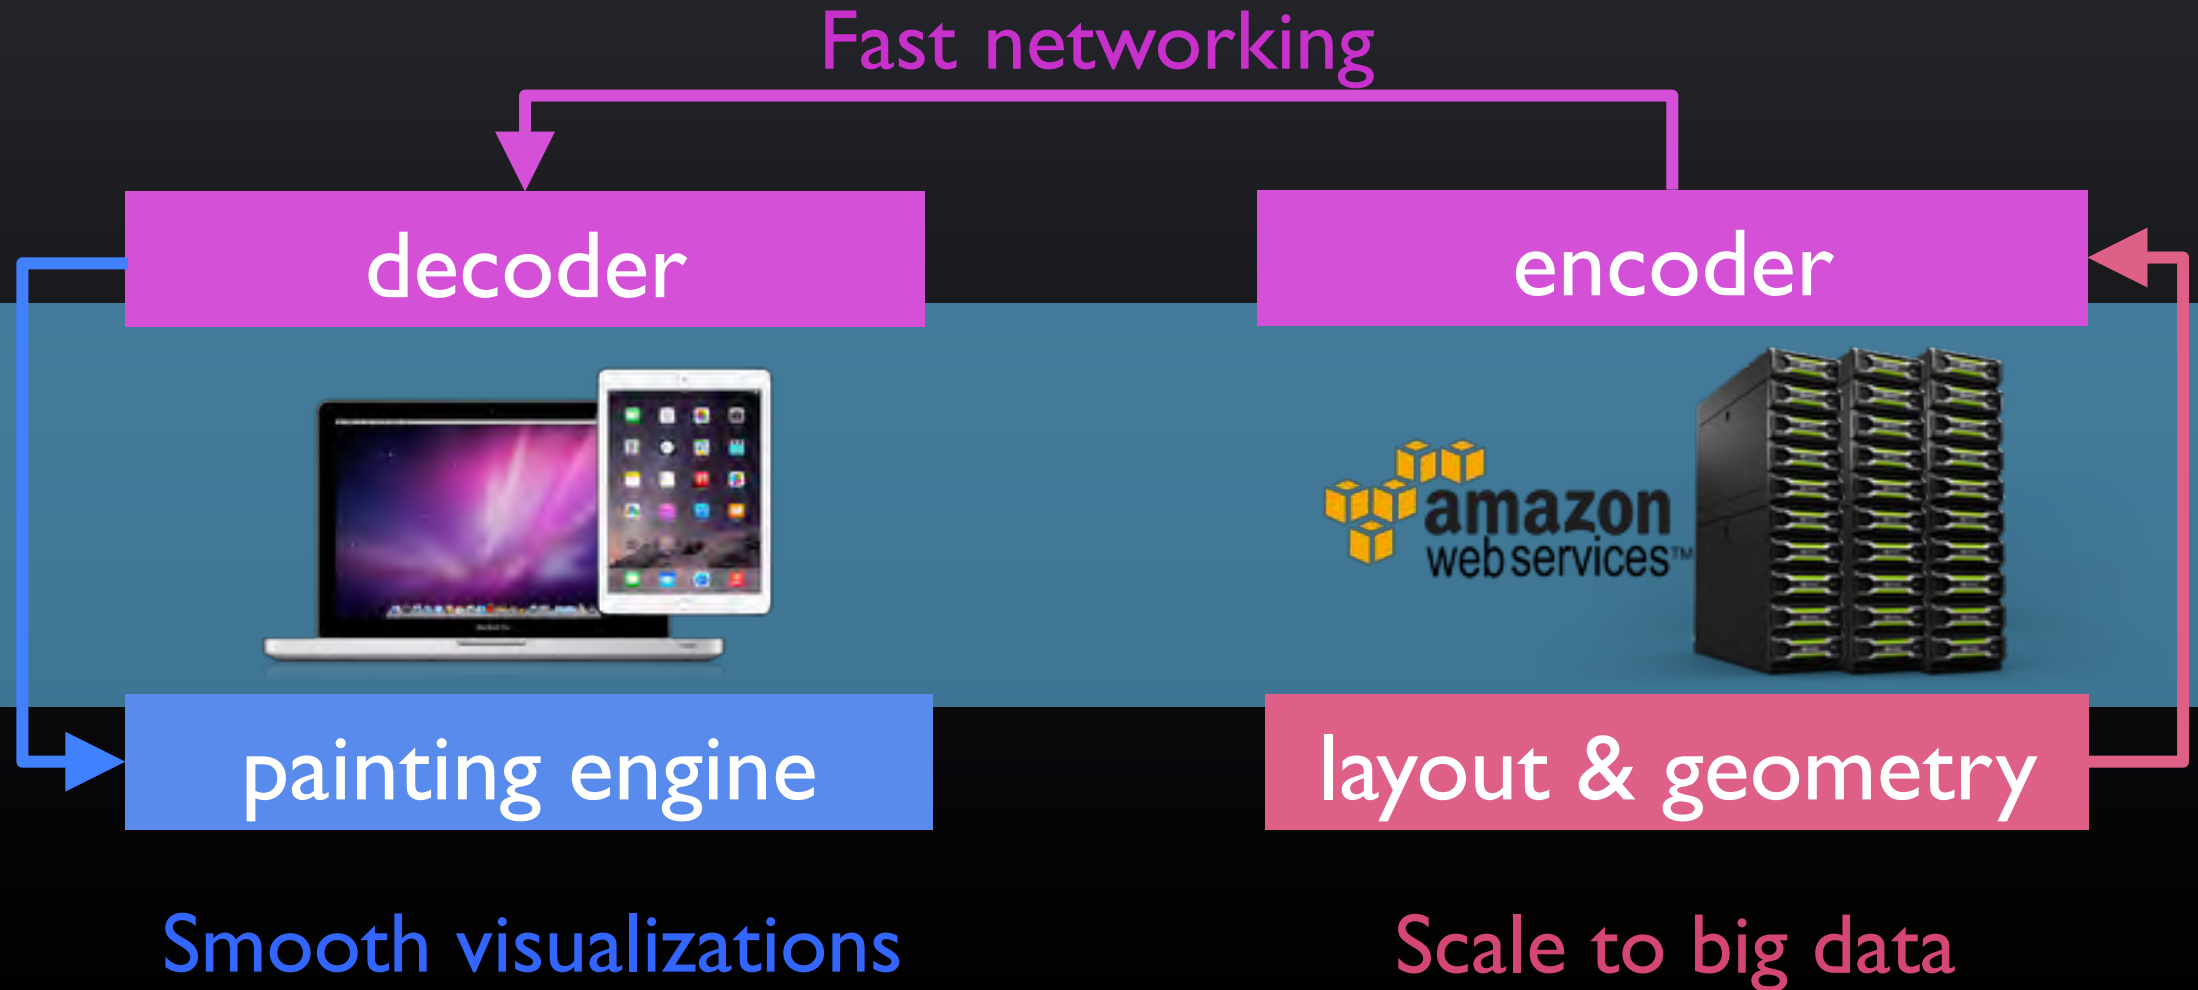
1. Distributed Rendering: Dual GPUs

2. GPU Programming with Node

General Programming with GPUs

Using node-opengl and cl.js for easy acceleration

Dual GPU Architecture



Why Not All Clientside?

We have WebGL, but...

1 GB+/minute input data

Client provides very restricted access to GPU features

Not even WebGL2 provides core performance features, e.g., memory barriers.

Why Not All Serverside?



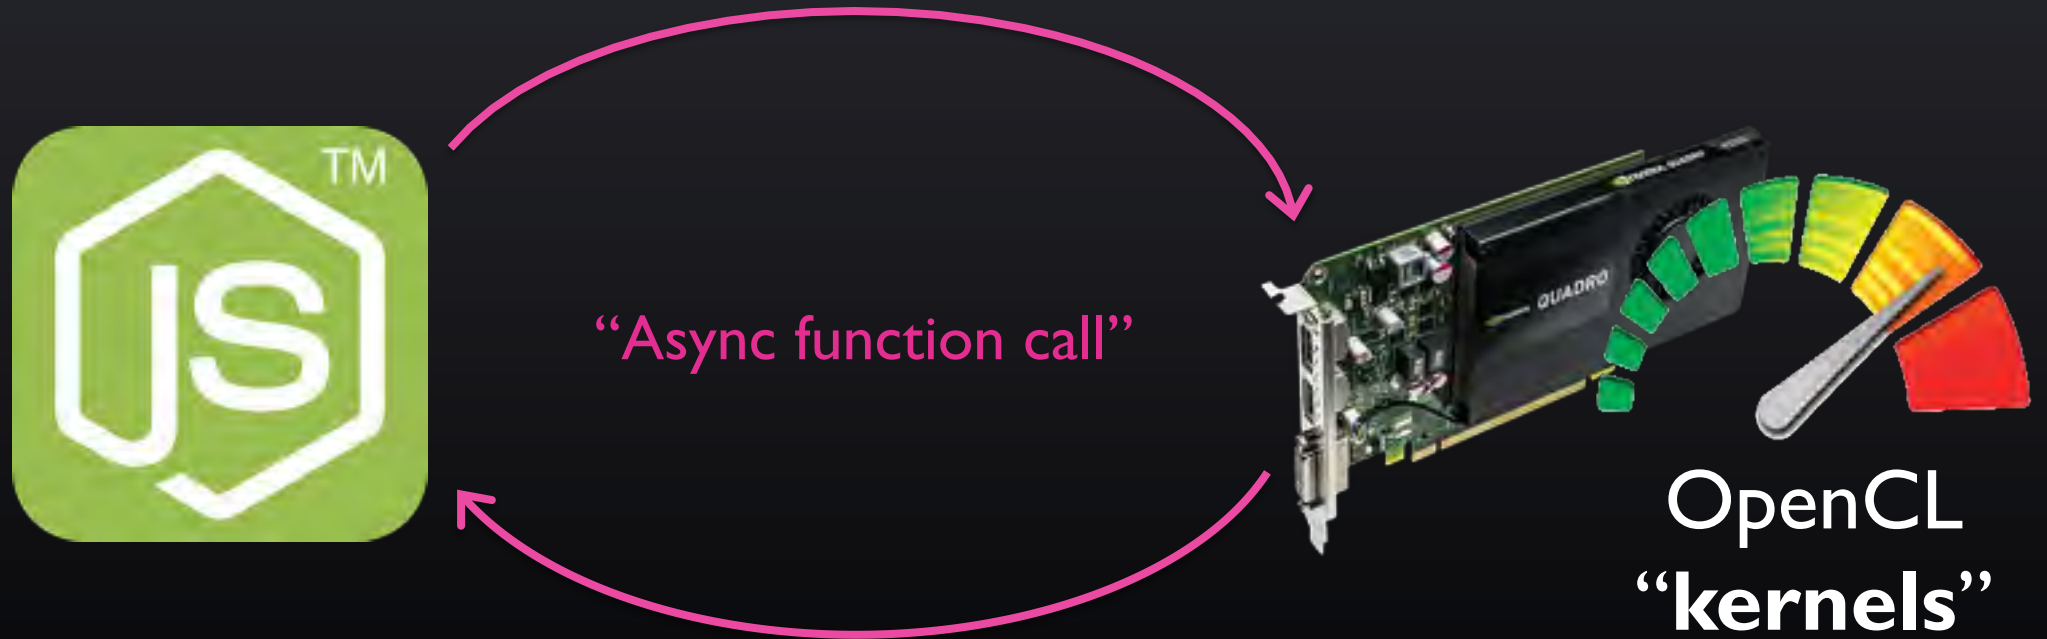


Low-level
CAPIs



See other
great talks!

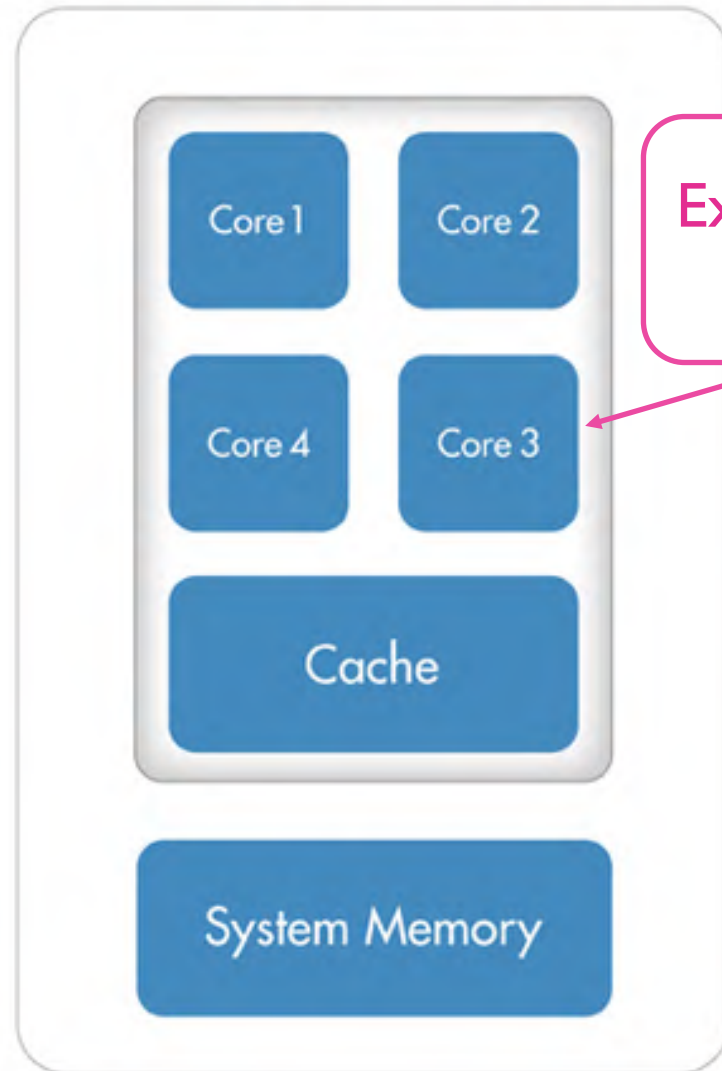
Always Bet on JS



JS community has (inadvertently) made GPU programming awesome!

CPU

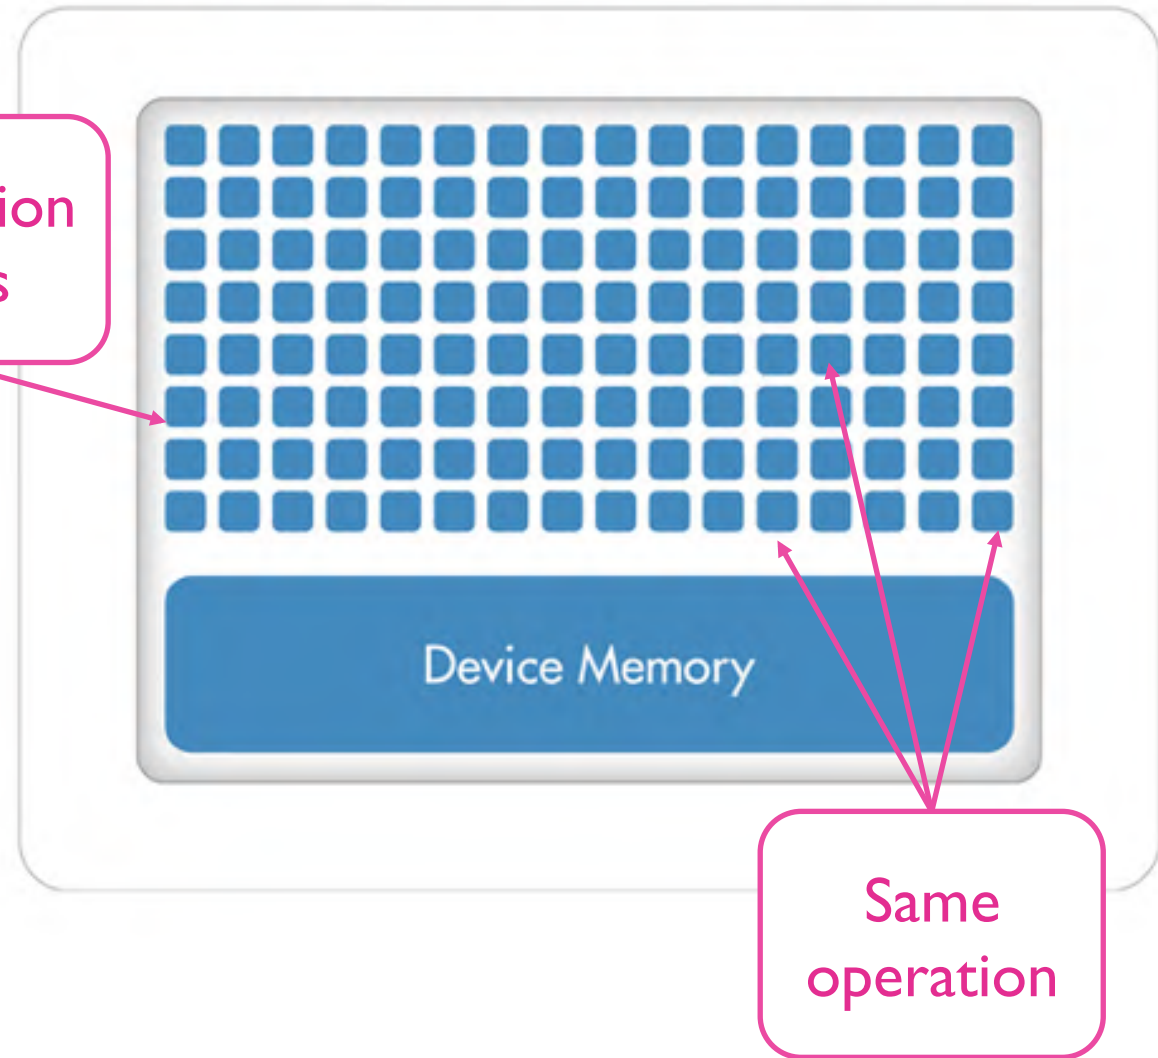
Few but Flexible



Execution
units

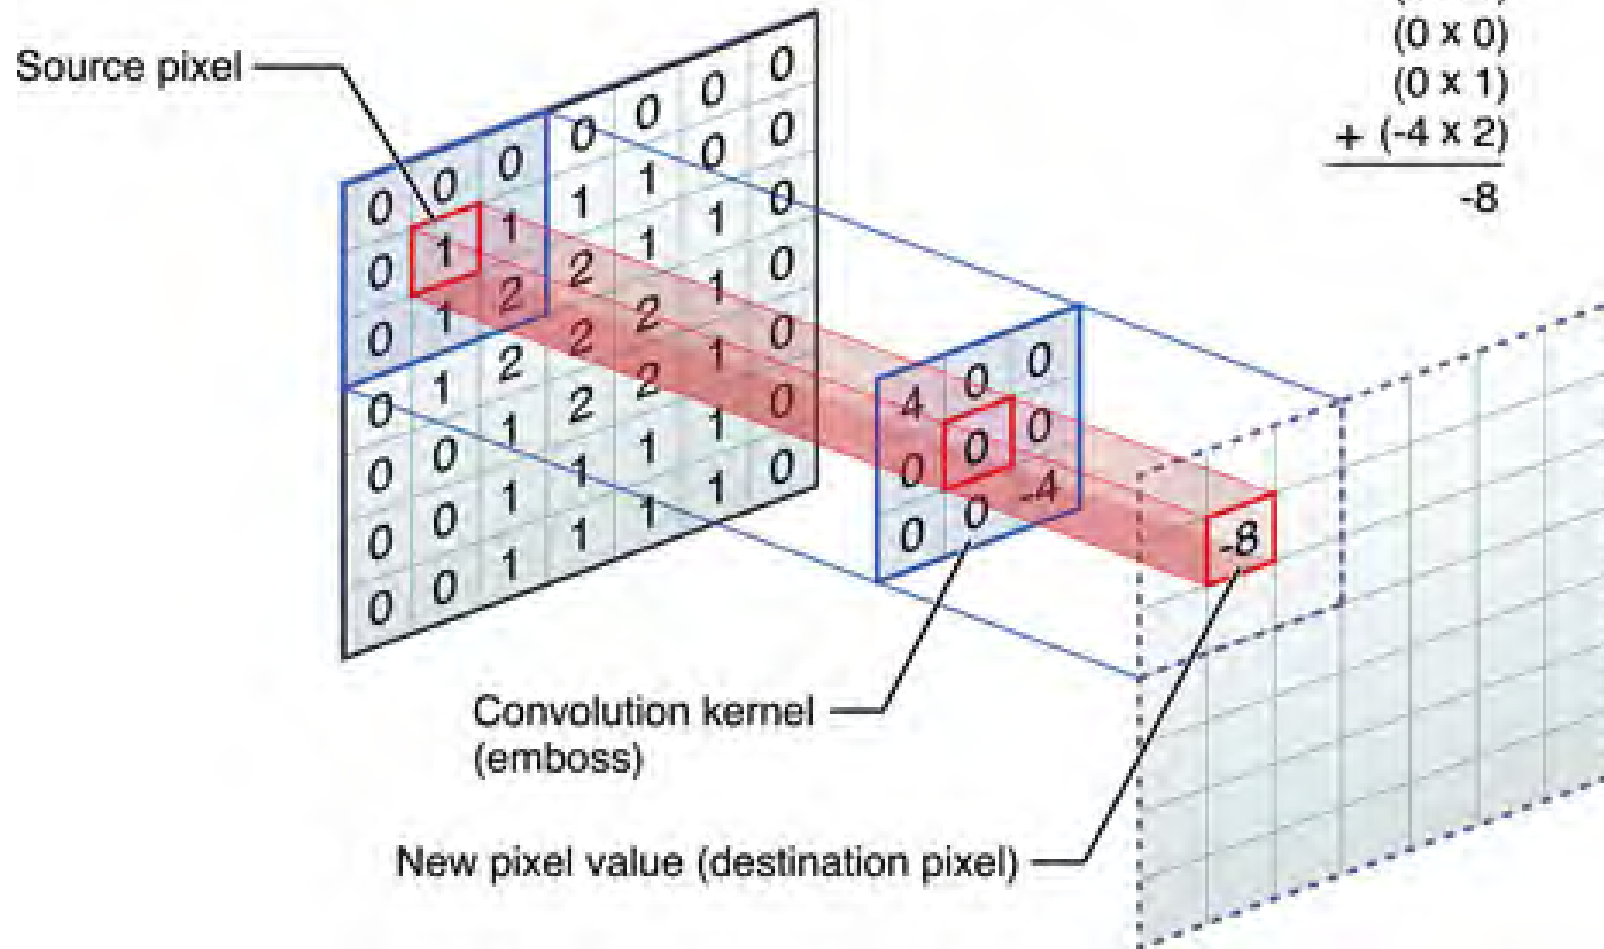
GPU

Many but Rigid



Same
operation

Center element of the kernel is placed over the source pixel. The source pixel is then replaced with a weighted sum of itself and nearby pixels.



$$\begin{array}{r}
 (4 \times 0) \\
 (0 \times 0) \\
 (0 \times 0) \\
 (0 \times 0) \\
 (0 \times 1) \\
 (0 \times 1) \\
 (0 \times 0) \\
 (0 \times 1) \\
 + (-4 \times 2) \\
 \hline
 -8
 \end{array}$$



```
int centerIdx = get_global_id(0);  
int sum = 0;
```

ID (Pixel)

```
for (int dx = -1; dx <= 1; dx++) {  
    for (int dy = -1; dy <= 1; dy++) {  
        int idx = centerIdx + (4 * width * dy) + (4 * dx);  
        idx = isValidImageIndex(idx, numElements) ? idx : centerIdx;  
        int scaleFromMask = mask[(dy+1)*3 + (dx+1)];  
        sum += ((int) imageData[idx]) * scaleFromMask;  
    }  
}
```

For Each Neighbor

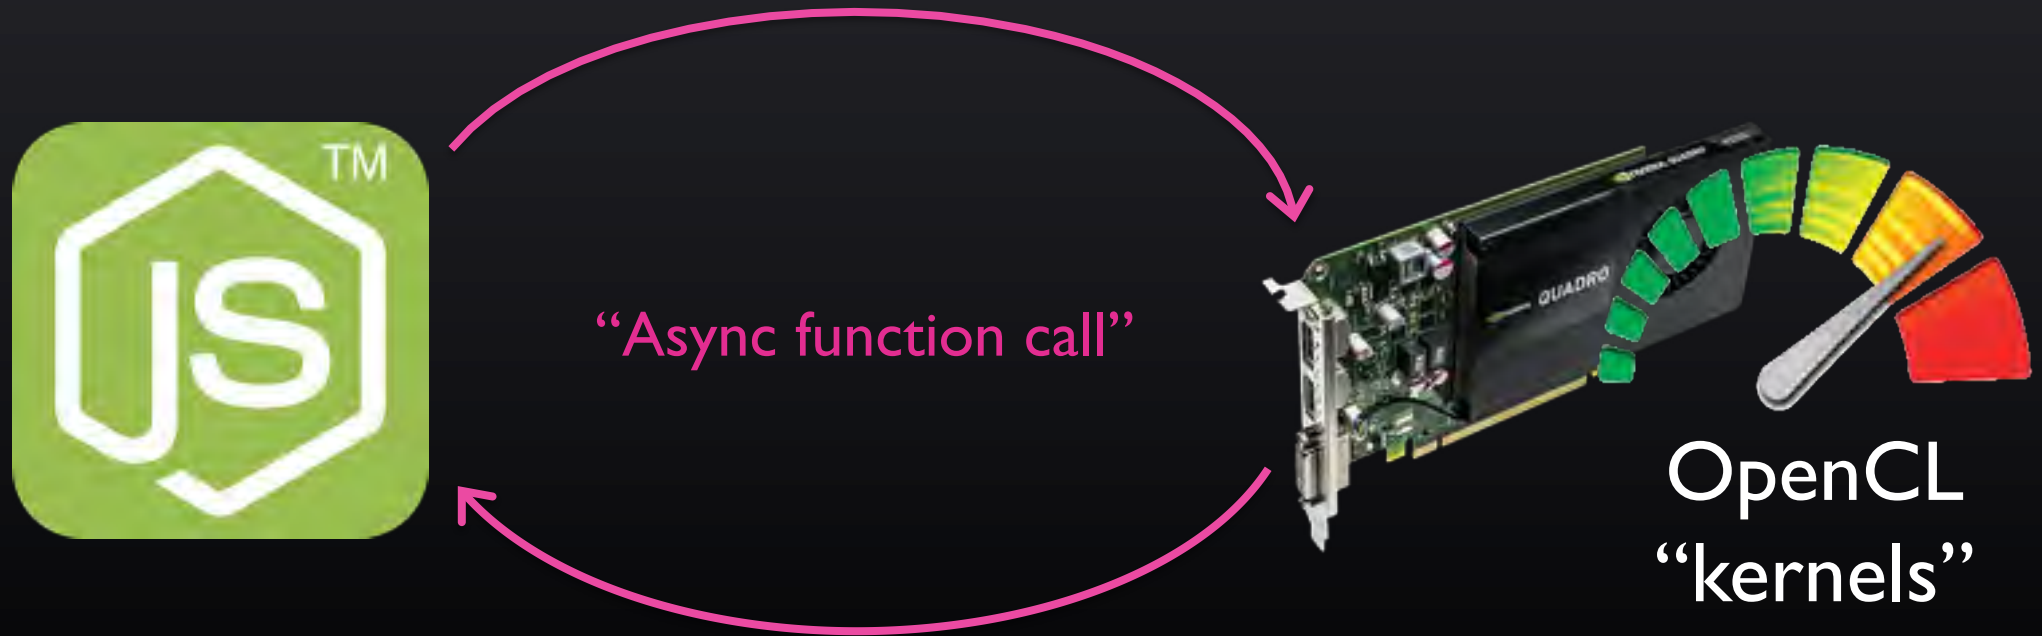
Do Math

```
if (isTransparencyNotColor(centerIdx)) {  
    sum = (uchar) 255;  
}
```


```
newImageData[centerIdx] = make8Bit(sum);
```

Write Result

Using Kernels From JavaScript



Node-openc1



Works on Intel,
Nvidia, and
AMD hardware

Developed by Mikael Sevenier at AMD.

Provides bindings in node to C++ driver functions.

Handles memory management, concurrency, etc., for you!


```
// Setup Context
var ocl = require('node-openccl');
var platforms = ocl.getPlatformIDs();
var platform = platforms[0];
var devices = ocl.getDeviceIDs(platform, ocl.DEVICE_TYPE_ALL);
var clErrorHandler = function (e) { throw e; };
var context = ocl.createContext([ocl.CONTEXT_PLATFORM, platform],
    device, clErrorHandler, clErrorHandler);
var queue = ocl.createCommandQueue(context, device, 0);
```

CL.js

```
var cl = require('cljs');

var myKernel = cl.createKernel('kernel.cl', 'kernelName');

var inputBuffer = cl.createBuffer(input);
var outputBuffer = cl.createBuffer(output);

return myKernel.run(
  [256],
  workgroup,
  [inputBuffer, outputBuffer]
).then(function () {
  return outputBuffer.read(Uint8Array);
});
```

Make a
kernel

Upload data
to GPU

Number of items (eg, pixels)

Kernel arguments

Download result

I Will Do It Live!

Takeaways



Use the very best tool for each task!



OpenCL

Performance (60x with one GPU)



Productivity (95% of code)

Thanks!

Try out

CL.js

github.com/graphistry/cljs

PyGraphistry

github.com/graphistry/pygraphistry

Help us
catch
hackers

 **graphistry**

We are hiring
Front-end / Full-stack
UI/UX design

Email us at
thibaud@graphistry.com

